

CS416: Systems Programming

Process Related Systems Call
(System, Fork, EXEC)

Creating processes

- **Method (1) using *system***

- The system function in the standard C library provides an easy way to execute a command from within a program, much as if the command had been typed into a shell.

```
#include <stdlib.h>
```

```
int main ()
```

```
{
```

```
    int return_value;
```

```
    return_value = system ("ls -l /");
```

```
    return return_value;
```

```
}
```

Creating processes

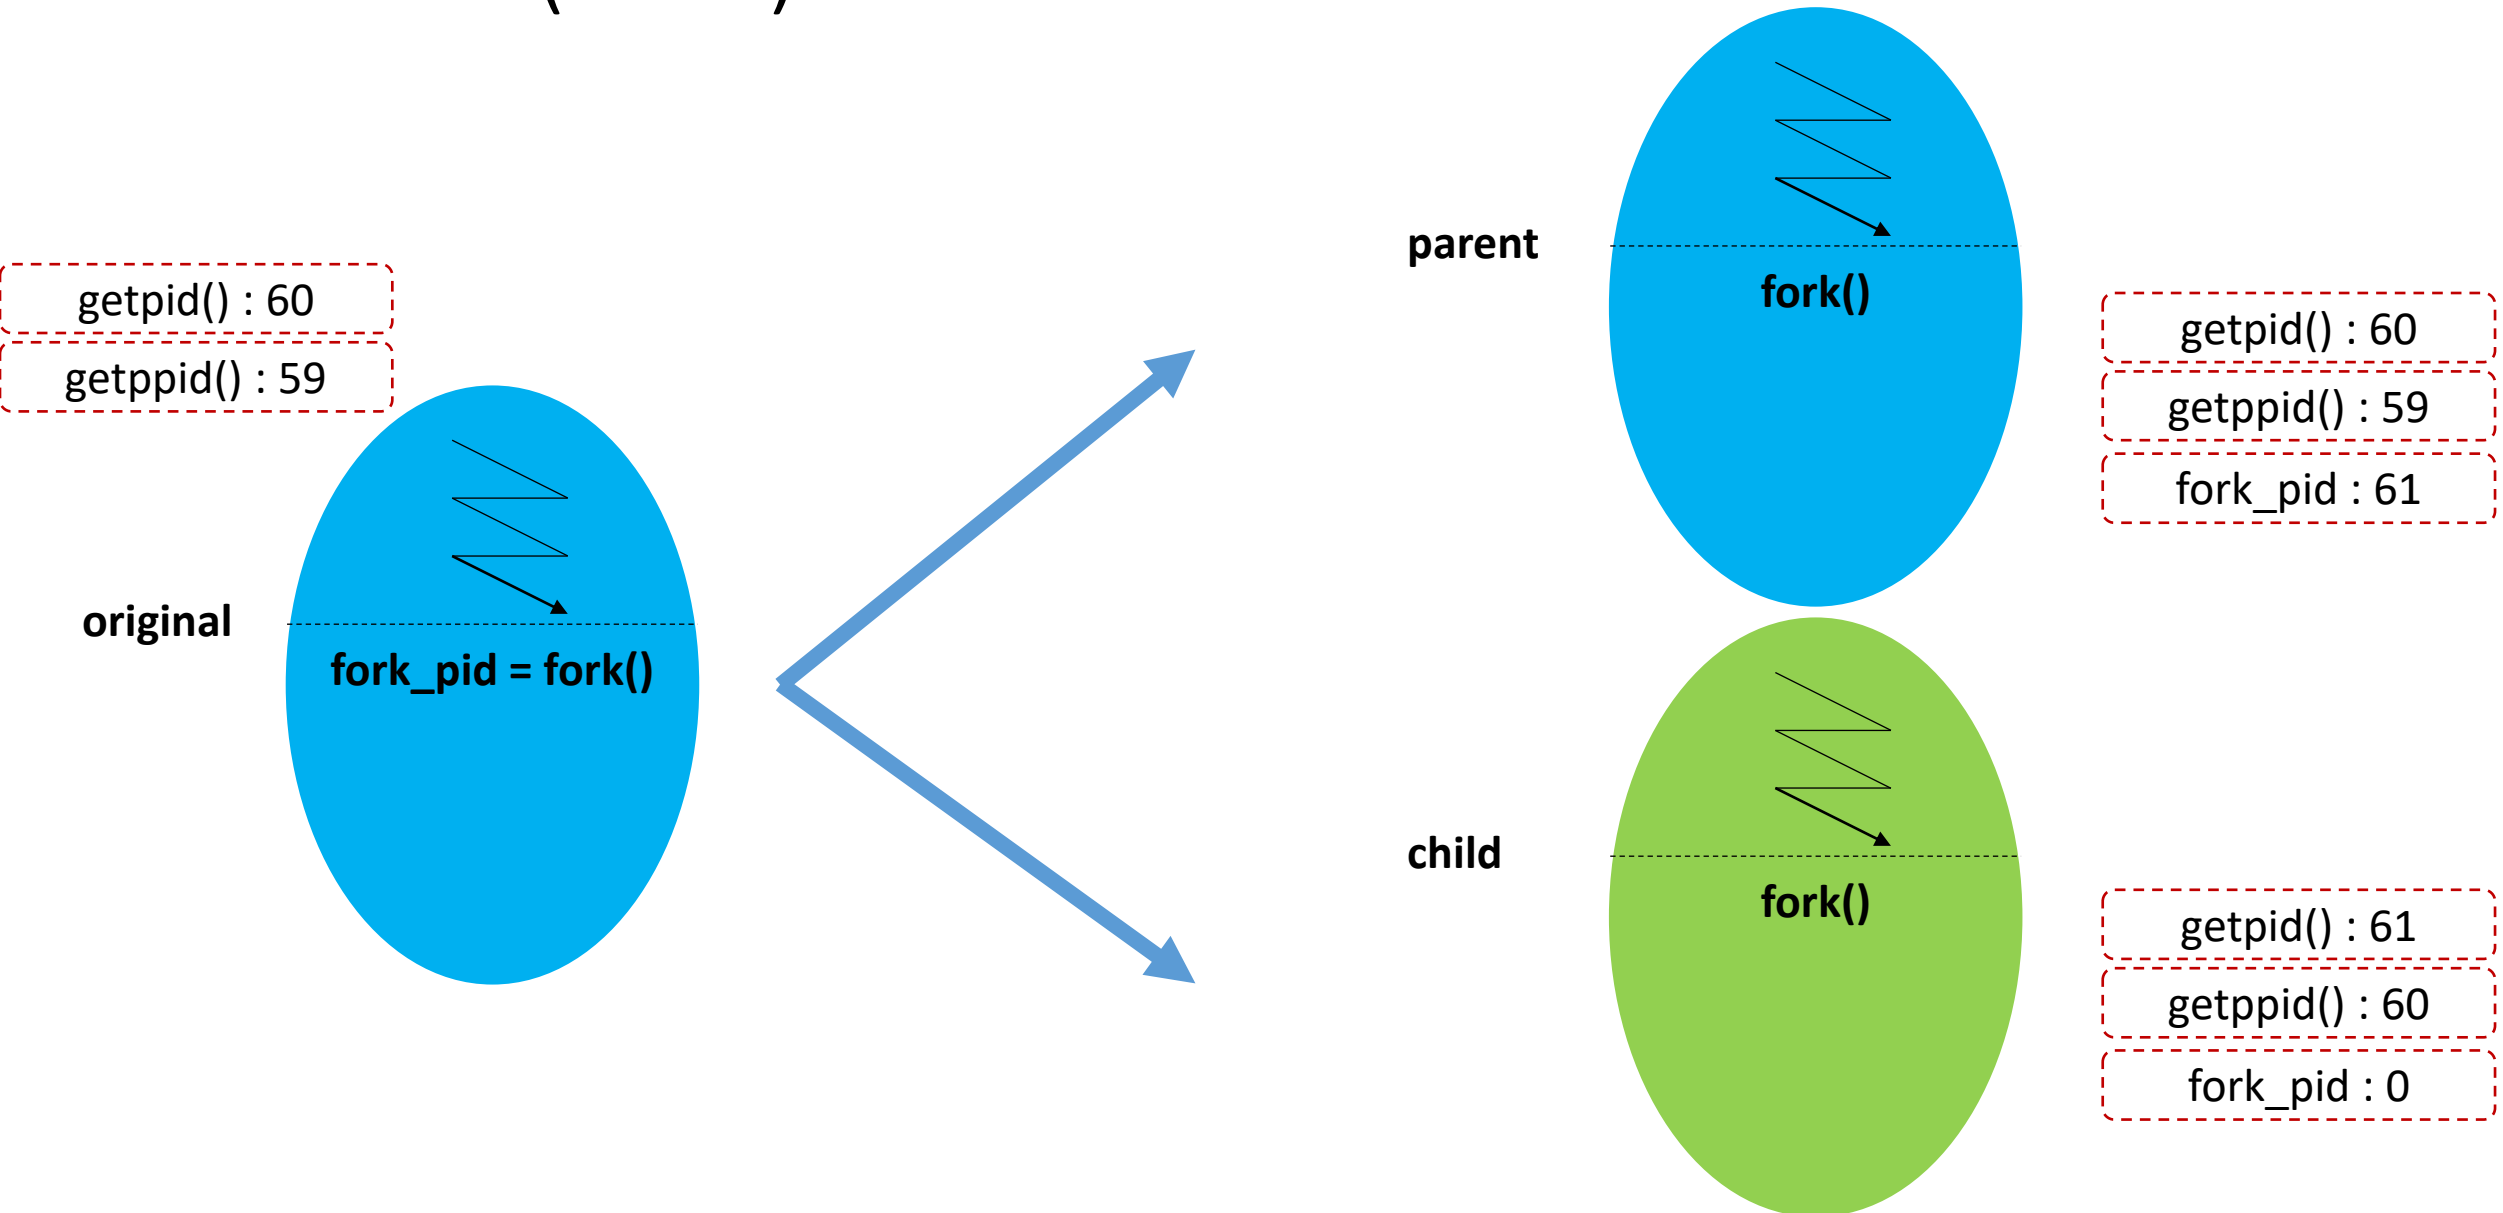
- **Method (2) using *fork* and *exec***

- Linux provides one function, **fork**, that makes a child process that is an exact copy of its parent process
- Linux provides another set of functions, the **exec** family, that causes a particular process to cease being an instance of one program and to instead become an instance of another program
- To **spawn** a new process, you first use **fork** to make a copy of the current process. Then you use **exec** to transform one of these processes into an instance of the program you want to spawn

int fork(void)

- **Description:** create a child process
- **Returns:** process ID of the new process

int fork(void)



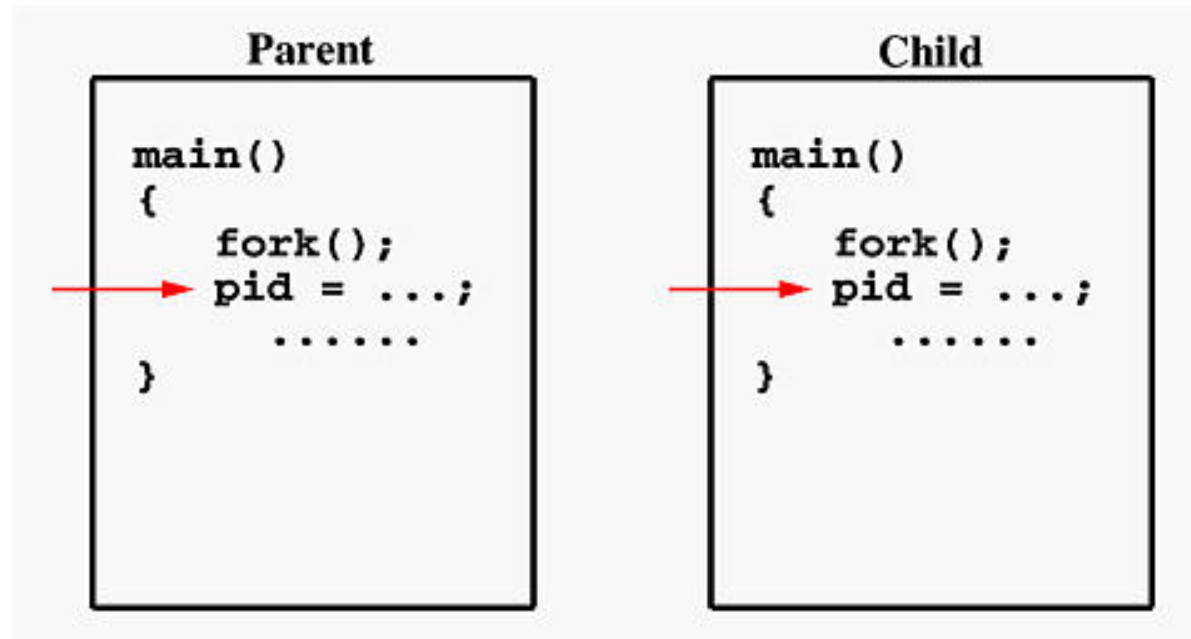
fork() - Example

```
1  #include <stdio.h>
2  #include <sys/types.h>
3  #include <sys/wait.h>
4  #include <unistd.h>
5
6  int main (void)
7  {
8      pid_t fork_pid;
9      printf ("the main program process ID is %d\n", (int) getpid ());
10     printf ("the main program parent process ID is %d\n", (int) getppid ());
11     fork_pid = fork ();
12
13     if (fork_pid != 0)
14     {
15         printf ("*****Parent Process*****\n");
16         printf ("process ID is %d\n", (int) getpid ());
17         printf ("parent process ID is %d\n", (int) getppid ());
18         printf ("the child's process ID is %d\n", (int) fork_pid);
19         sleep(10);
20     }
21     else{
22         wait(NULL);
23         printf ("*****Child Process*****\n");
24         printf ("process ID is %d\n", (int) getpid ());
25         printf ("parent process ID is %d\n", (int) getppid ());
26         printf ("logical ID of the process based on the fork function is %d\n", (int) fork_pid);
27     }
28     return 0;
29 }
```

```
➤ ./main
the main program process ID is 117
the main program parent process ID is 6
*****Parent Process*****
process ID is 117
parent process ID is 6
the child's process ID is 118
*****Child Process*****
process ID is 118
parent process ID is 117
logical ID of the process based on the fork function is 0
```

Contd.

- System call **fork()** is used to create processes.
- It takes no arguments and returns a process ID.
- The purpose of **fork()** is to create a **new** process, which becomes the *child* process of the caller.
- After a new child process is created, **both** processes will execute the next instruction following the **fork()** system call.



Distinguish parent from child processes

Simply check the returned value of **fork()**:

- If **fork()** returns a negative value, the creation of a child process was unsuccessful.
- **fork()** returns a zero to the newly created child process.
- **fork()** returns a positive value, the ***process ID*** of the child process, to the parent.

The returned process ID is of type **pid_t** defined in **sys/types.h**.

Normally, the process ID is an integer. Moreover, a process can use function **getpid()** to retrieve the process ID assigned to this process.

Another example

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#define MAX_COUNT 3
void ChildProcess(void); /* child process prototype */
void ParentProcess(void); /* parent process prototype */
```

```
int main(void)
{
    pid_t pid;
    pid = fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
    return 0;
}

void ChildProcess(void)
{
    int i;
    for (i = 1; i <= MAX_COUNT; i++)
        printf(" This line is from child, value = %d\n", i);
    printf(" *** Child process is done ***\n");
}

void ParentProcess(void)
{
    int i;
    for (i = 1; i <= MAX_COUNT; i++)
        printf("This line is from parent, value = %d\n", i);
    printf("*** Parent is done ***\n");
}
```

Previous Example - Sample Output

```
This line is from parent, value = 1
This line is from parent, value = 2
This line is from parent, value = 3
*** Parent is done ***
    This line is from child, value = 1
    This line is from child, value = 2
    This line is from child, value = 3
*** Child process is done ***
```

```
This line is from parent, value = 1
    This line is from child, value = 1
    This line is from child, value = 2
    This line is from child, value = 3
    *** Child process is done ***
This line is from parent, value = 2
This line is from parent, value = 3
*** Parent is done ***
```

```
./main
This line is from parent, value = 1
This line is from parent, value = 2
    This line is from child, value = 1
    This line is from child, value = 2
    This line is from child, value = 3
    *** Child process is done ***
This line is from parent, value = 3
*** Parent is done ***
```

```
This line is from parent, value = 1
This line is from parent, value = 2
This line is from parent, value = 3
*** Parent is done ***
```

How does it work

Parent

Child

```
main()
{
    pid = 3456
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```

```
main()
{
    pid = 0
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```

The execution depends on the fork
returned values

Till this point there is no difference

Parent

Child

```
main()
{
    pid = 3456
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

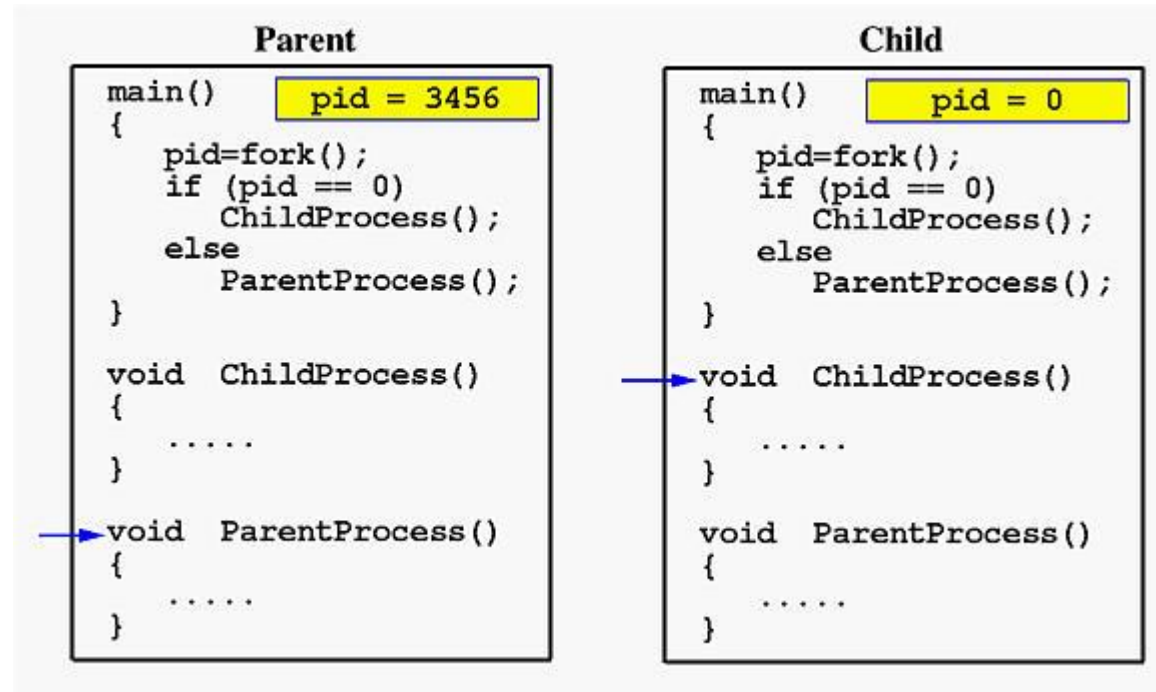
void ParentProcess()
{
    .....
}
```

```
main()
{
    pid = 0
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```

How does it work



Based on the fork returned value, one function is chosen. If both functions have an output command (like `printf`), the outputs in the child and parent processes might intersect

One extra example

```
#include <stdio.h>
#include <unistd.h>
```

```
int main(int argc, char **argv)
{
    printf("--beginning of program\n");

    int counter = 0;
    pid_t pid = fork();

    if (pid == 0)
    {
        // child process
        int i = 0;
        for (; i < 5; ++i)
        {
            printf("child process: counter=%d\n", ++counter);
        }
    }
    else if (pid > 0)
    {
        // parent process
        int j = 0;
        for (; j < 5; ++j)
        {
            printf("parent process: counter=%d\n", ++counter);
        }
    }
    else
    {
        // fork failed
        printf("fork() failed!\n");
        return 1;
    }
    printf("--end of program--\n");
    return 0;
}
```

```
--beginning of program
parent process: counter=1
parent process: counter=2
parent process: counter=3
parent process: counter=4
parent process: counter=5
--end of program--
child process: counter=1
child process: counter=2
child process: counter=3
child process: counter=4
child process: counter=5
--end of program--
```

Change the program running in the process

The **exec**..... function replaces the program running in a process with another program.

The exec family has many similar functions that vary in the way they are called

1. Functions with the letter **p** in their names (execvp and execlp) accept a program name and search for it in the current execution path. No p in the name means you need to provide the full path of the program.
2. Functions with the letter **v** in their names (execv, execvp, and execve) accept an argument list for the program as a null terminated array of pointers to strings
3. Functions with the letter **e** in their names (execve and execl e) accept an additional argument, an array of environment variables.

More about exec

- Exec family never returns unless an error occurs
 - Use the function `fprintf`
 - Need the library `stdlib.h`
- Programmer should pass the name of the program as the first element of the argument list


```
int execv(char *progName, char *argv[])
```

- description: load program in process address space
- returns: -1 if failed




execv()


original




parent



child



New program



prog1.c

```
#include <stdio.h>
#include <unistd.h>

int main(void) {
    printf("Child Process: Hello -- %d\n", getpid());
    return 0;
}
```

```
>> gcc prog1.c -o childProg
```

```
>> ./childProg
```

```
~/WildMammothMicroscope$ gcc prog1.c -o childProg
~/WildMammothMicroscope$ ./childProg
Child Process: Hello -- 1014
~/WildMammothMicroscope$
```

prog2.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <stdlib.h>
int main()
{
    char* arg_list[]={"/.childProg", NULL};
    pid_t child_pid;
    child_pid=fork();
    if(child_pid!=0)
    {
        printf("Parent Process: Hello -- %d\n", getpid());
        sleep(5); //sleep for 5 seconds
    }
    else{
        printf("Child Process: Before being replaced with another program -- %d\n", getpid());
        execvp(arg_list[0],arg_list);
        printf("Child Process: After being replaced with another program\n");
        abort();
    }
    printf("Done with the main prgram\n");
    return 0;
}
```

```
>> gcc prog2.c -o parentProg
```

```
>> ./parentProg
```

```
~/WildMammothMicroscope$ gcc prog2.c -o parentProg
~/WildMammothMicroscope$ ./parentProg
Parent Process: Hello -- 1528
Child Process: Before being replaced with another program -- 1529
Child Process: Hello -- 1529
Done with the main prgram
~/WildMammothMicroscope$
```

prog1.c

```
#include <stdio.h>
#include <unistd.h>

int main(void) {
    printf("Child Process: Hello -- %d\n", getpid());
    return 0;
}
```

```
>> gcc prog1.c -o childProg
```

```
>> ./childProg
```

```
~/WildMammothMicroscope$ gcc prog1.c -o childProg
~/WildMammothMicroscope$ ./childProg
Child Process: Hello -- 1014
~/WildMammothMicroscope$
```

prog2.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <stdlib.h>
int main()
{
    char* arg_list[]={"/NotExistingProg", NULL};
    pid_t child_pid;
    child_pid=fork();
    if(child_pid!=0)
    {
        printf("Parent Process: Hello -- %d\n", getpid());
        sleep(5); //sleep for 5 seconds
    }
    else{
        printf("Child Process: Before being replaced with another program -- %d\n", getpid());
        execvp(arg_list[0],arg_list);
        printf("Child Process: After being replaced with another program\n");
        abort();
    }
    printf("Done with the main prgram\n");
    return 0;
}
```

```
>> gcc prog2.c -o parentProg
```

```
>> ./parentProg
```

```
~/WildMammothMicroscope$ gcc prog2.c -o parentProg
~/WildMammothMicroscope$ ./parentProg
Parent Process: Hello -- 2161
Child Process: Before being replaced with another program -- 2162
Child Process: After being replaced with another program
Done with the main prgram
~/WildMammothMicroscope$
```

fork and exec (Create spawn function for Linux)

```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <sys/types.h>
void spawn(char* program, char** arg_list)
{
    pid_t child_pid;
    child_pid=fork();
    if(child_pid!=0)
    {
        printf("Run\n");
        printf("    Diff.\n");
        printf("    Things\n");
        sleep(5); //sleep for 5 seconds
    }
    else{
        execvp(program,arg_list);
        fprintf(stderr,"an error occured in execvp\n");
        abort();
    }
}

int main()
{
    char* arg_list[]={ "ls", "-l", "/", NULL};
    spawn(arg_list[0], arg_list);

    printf("Done with the main prgram\n");
    return 0;
}
```

- The Linux **"ls"** command allows you to view a list of the files and folders in a given directory.
- The **-l** flag instructs Linux to print out a list of files with detailed descriptions.
- A forward slash **"/"** is used to refer to the root directory.

Correct and false runs

```
❖ ./main
Run
Diff.
Things
total 88
drwxr-xr-x 1 root root 4096 Feb 29 21:32 bin
drwxr-xr-x 2 root root 4096 Apr 24 2018 boot
drwxr-xr-x 1 runner runner 4096 Dec 4 19:14 config
drwxr-xr-x 5 root root 340 Feb 29 21:32 dev
drwxr-xr-x 1 root root 4096 Feb 29 21:32 etc
drwxr-xr-x 3 root root 4096 Nov 21 00:03 hom
drwxr-xr-x 1 root root 4096 Nov 21 00:01 home
dr-xr-xr-x 4 nobody nogroup 4096 Feb 29 21:32 io
drwxr-xr-x 1 root root 4096 Nov 21 00:01 lib
drwxr-xr-x 2 root root 4096 Nov 20 23:50 lib32
drwxr-xr-x 2 root root 4096 Oct 29 21:25 lib64
drwxr-xr-x 2 root root 4096 Oct 29 21:25 media
drwxr-xr-x 2 root root 4096 Oct 29 21:25 mnt
drwxr-xr-x 1 root root 4096 Dec 4 19:14 opt
dr-xr-xr-x 777 nobody nogroup 0 Feb 29 21:32 proc
drwx----- 1 root root 4096 Dec 4 19:13 root
drwxr-xr-x 1 root root 4096 Dec 4 19:13 run
drwxr-xr-x 1 root root 4096 Dec 4 19:14 run_dir
drwxr-xr-x 1 root root 4096 Feb 29 21:32 sbin
drwxr-xr-x 2 root root 4096 Oct 29 21:25 srv
dr-xr-xr-x 13 nobody nogroup 0 Feb 29 10:25 sys
drwxrwxrwt 1 root root 4096 Feb 29 22:46 tmp
drwxr-xr-x 1 root root 4096 Nov 21 00:15 usr
drwxr-xr-x 1 root root 4096 Nov 21 00:01 var
Done with the main program
❖
```

```
❖ ./main
Run
Diff.
Things
an error occured in execvp
Done with the main program
❖
```

execvp() vs. execl()

```
#include <stdio.h>

#include <unistd.h>

#include <stdlib.h>

#include <sys/types.h>

int main()

{

    pid_t child_pid;

    child_pid=fork();

    if(child_pid!=0){

        printf("Run\n");

        printf(" Diff.\n");

        printf(" Things\n");

        sleep(5); //sleep for 5 seconds

    } else{

        char* arg_list[]={ "ls", "-l", "/", NULL};

        execvp(arg_list[0],arg_list);

        fprintf(stderr,"an error occured in execvp\n");

        abort();

    }

    printf("Done with the main prgram\n");

    return 0;

}
```

```
#include <stdio.h>

#include <unistd.h>

#include <stdlib.h>

#include <sys/types.h>

int main()

{

    pid_t child_pid;

    child_pid=fork();

    if(child_pid!=0){

        printf("Run\n");

        printf(" Diff.\n");

        printf(" Things\n");

        sleep(5); //sleep for 5 seconds

    } else{

        execl("/bin/ls","ls", "-l", "/", NULL);

        fprintf(stderr,"an error occured in execvp\n");

        abort();

    }

    printf("Done with the main prgram\n");

    return 0;

}
```

Linux vs. Windows

- The DOS and windows API contain the *spawn* family of functions. This family takes the name of a program and creates a new process instance of it
- Linux does not contain a single function to perform the same task.
- Linux uses *fork* and *exec*.
 - *fork* makes the new process
 - *exec* transforms the process instance into the needed program

```
int spawnl(int mode, char *path, char *arg0, ...);  
int spawnle(int mode, char *path, char *arg0, ..., char ** envp);  
int spawnlp(int mode, char *path, char *arg0, ...);  
int spawnlpe(int mode, char *path, char *arg0, ..., char ** envp);  
int spawnv(int mode, char *path, char **argv);  
int spawnve(int mode, char *path, char **argv, char ** envp);  
int spawnvp(int mode, char *path, char **argv);  
int spawnvpe(int mode, char *path, char **argv, char ** envp);
```

Name	Notes
e	An array of pointers to environment arguments is explicitly passed to the child process.
l	Command line arguments are passed individually to the function.
p	Uses the PATH argument variable to find the file to be executed.
v	Command line arguments are passed to the function as an array of pointers.